

Ridge Regression on Diamonds Dataset

A Comparative Analysis of Optimization Algorithms

Davide Gamba

Matr. 1053470

University of Bergamo

Department of Management, Information and Production Engineering

March 30, 2026



Introduction

The objective of this study is to conduct a comparative analysis of the convergence behavior of various optimization algorithms applied on a Ridge Regression problem to the Diamonds Dataset.

In particular, the focus is on:

- Evaluating the respective algorithms convergence rates.
- Investigating how theoretical properties of the objective function (such as Convexity, Smoothness ecc..) have an impact to results.
- Analyzing how the L_2 regularization parameter (λ) of Ridge Regression influences algorithms behavior.

The following algorithms were implemented:

- Gradient Descent (Naive, Bounded, Smooth).
- Accelerated Gradient Descent (Nesterov's).
- Mini-Batch Stochastic Gradient Descent (Decaying Learning rate)
- Newton's Method.
- Adagrad.

Dataset Overview: The Diamonds Dataset

We utilized the **Diamonds Dataset**. The dataset consists of approximately **54,000 observations** ($n = 53.940$), each representing a diamond with **10 distinct attributes**. The goal is to predict with a Ridge Regression the **Price** of a diamond based on its physical and quality characteristics.

Feature Description:


- **Target Variable (y):** price in US dollars.
- **Numerical Features:**
 - carat: Weight of the diamond.
 - x, y, z: Length, Width, and Depth in mm.
 - depth: The depth of the diamond divided by the diameter of the diamond.
 - table: The length of the table facet divided by the diameter of the diamond.
- **Categorical Features:**
 - cut: Quality of the cut (Fair, Good, Very Good, Premium, Ideal).
 - color: Diamond color, from J (coloured, worst) to D (colorless, best).
 - clarity: How clear the material is and what surface imperfections it have.

Dataset Overview: The Diamonds Dataset

Distribution of all the main variables within the diamonds dataset.



Dataset Characteristics: Feature matrix



Dataset Characteristics: Feature correlation

Multicollinearity:

In our dataset, there is a strong correlation between all variables, in particular between carat (weight) and the physical dimensions (x , y , z).

Impact on Optimization:

- Given the high multicollinearity among the features in the matrix X , we expect the loss objective landscape exhibit flat and elongated valleys.
- Consequently, the optimization problem is **ill-conditioned** and particularly challenging for gradient-based methods.

Ridge Regression

Ridge Regression problem extends Ordinary Least Squares (OLS) by imposing an L_2 penalty on the model parameters to mitigate overfitting and multicollinearity.

The objective function is the mean squared errors for each data point ($i = 1$ to n) plus the sum of squared weights (from $j = 1$ to d).

$$f(w) = \frac{1}{2n} \sum_{i=1}^n (y_i - x_i^T w)^2 + \frac{\lambda}{2} \sum_{j=1}^d w_j^2$$

We can define it in vector notation as follows:

$$f(w) = \underbrace{\frac{1}{2n} \|y - Xw\|^2}_{\text{Half MSE}} + \underbrace{\frac{\lambda}{2} \|w\|^2}_{\text{Regularization}}$$

Where:

- $X \in \mathbb{R}^{n \times d}$ is the feature matrix (includes a column of 1).
- $y \in \mathbb{R}^n$ is the target variable (price).
- $w \in \mathbb{R}^d$ is the weights vector.
- $\lambda > 0$ is the regularization hyperparameter.

Gradient and Hessian Matrix

Since $f(w)$ is differentiable everywhere, we can compute the first-order derivative (the Gradient) and second-order derivative (the Hessian matrix).

The Gradient $\nabla f(w)$

$$\nabla f(w) = \frac{1}{n} X^T (Xw - y) + \lambda w$$

The Hessian Matrix $\nabla^2 f(w)$

$$H = \nabla^2 f(w) = \frac{1}{n} X^T X + \lambda I$$

$f(w)$ is strictly convex

To investigate the convexity of our objective function, we rely on the Second-order Characterization of Convexity.

A symmetric matrix M is Positive Semi-Definite if $v^T M v \geq 0$ for all vectors v , and Positive Definite if $v^T M v > 0$ for all non-zero vectors $v \neq 0$.

Let's apply this definition to our Ridge Regression Hessian matrix, analyzing its quadratic form $v^T H v$ for any arbitrary non-zero vector $v \in \mathbb{R}^d$:

$$v^T H v = v^T \left(\frac{1}{n} X^T X + \lambda I \right) v$$

$f(w)$ is strictly convex

Solving the quadratic form, we get:

$$v^T H v = \underbrace{\frac{1}{n} \|Xv\|^2}_{\text{OLS Base}} + \underbrace{\lambda \|v\|^2}_{\text{Ridge Penalty}}$$

- **The OLS Base:** The squared norm $\|Xv\|^2$ is a squared length, meaning it is always non-negative (≥ 0).
- **The Ridge Penalty:** For any non-zero vector $v \neq 0$, its squared norm $\|v\|^2$ is strictly positive. Since our regularization hyperparameter is strictly positive ($\lambda > 0$), the entire term $\lambda \|v\|^2$ is strictly greater than zero (> 0).

Because we are adding a strictly positive term to a non-negative term, the entire sum is guaranteed to be strictly positive.

This proves that the Ridge Hessian matrix H is **Positive Definite** and consequently the objective function possesses a single unique global minimum, so that means $f(w)$ is strictly convex.

Ridge Regression - Properties

To understand how optimization algorithms will behave on Ridge Regression objective function $f(w)$, we must analyze its properties.

1. Convexity: Yes (Strictly Convex)

As proven in the previous section, the addition of the L_2 penalty ensures that the Hessian matrix H is Positive Definite. Therefore, the function is not just convex, but **strictly convex**. This guarantees that any local minimum is the unique global minimum w^* .

2. Lipschitz Continuous: No (Globally)

A function $f(w)$ is globally B -Lipschitz if its gradients are bounded in magnitude everywhere ($\|\nabla f(w)\| \leq B$). Our objective function $f(w)$ is a quadratic function. As the weights w can grow toward infinity, the gradient $\nabla f(w) = \frac{1}{n}X^T(Xw - y) + \lambda w$ also can grow linearly toward infinity. Because the gradient is unbounded over the entire domain $\mathbb{R}^d$, the function itself is **not globally Lipschitz continuous**.

3. Smoothness (L -smooth): Yes

It is L -**Smooth**. The curvature does not change abruptly and the gradient doesn't oscillate wildly, the function always lies *below* a certain parabola:

$$f(y) \leq f(x) + \nabla f(w)^T (y - x) + \frac{L}{2} \|x - y\|^2 \quad \forall x, y$$

- **Value of L :** It corresponds to the largest eigenvalue of the Hessian matrix (the maximum curvature of the loss landscape).

$$L = \lambda_{\max} \left(\frac{1}{n} X^T X \right) + \lambda$$

- **Consequence:** This value acts as the absolute **speed limit** for Gradient Descent. The optimal safe step size is $\gamma = \frac{1}{L}$.

4. Strong Convexity (μ -strongly convex): Yes

It is μ -**Strongly Convex**. The curvature is "at least" μ in every direction. The function always lies *above* the quadratic function:

$$f(y) \geq f(x) + \nabla f(x)^T (y - x) + \frac{\mu}{2} \|x - y\|^2 \quad \forall x, y$$

- **Value of μ :** It corresponds to the smallest eigenvalue of the Hessian matrix.


$$\mu = \lambda_{\min} \left(\frac{1}{n} X^T X \right) + \lambda$$

- **Consequence:** μ indicates the *minimum curvature* of the loss landscape.

The effect of $L2$ regularization

Specifically, increasing λ directly increases μ , the **strong convexity parameter**. When we set a large λ the minimum curvature (μ) is strictly bounded away from zero, the loss function becomes steeper and no longer has any flat regions where Gradient's based method might stall.

So we expect that optimization algorithms converge fastly on problems where regularization is stronger.



Dataset pre-processing steps

To prepare the dataset for our optimization algorithms, we now apply the following preprocessing steps:

- **Ordinal Encoding:** Categorical variables (cut, color, clarity) were mapped to integer values, preserving their inherent hierarchy.
- **Removing Outliers:** We filter out extreme anomalous data points.
- **Standardize the Features (X) and the Target (y):** By subtracting the mean and dividing by the standard deviation, we ensure all features and the target share the same scale (a mean of 0 and a variance of 1), also it allows Gradient Descent to point much more directly toward the minimum.
- **The Bias Trick (Adding a column of 1s):** We append a column of ones to our feature matrix X to absorb the intercept (or bias) term directly into our weight vector w .

Testing approach

To test the performance of the optimization algorithms, we use the following framework:

1. The Methodology:

- We define a specific set of λ (regularization) values:
 $\lambda \in \{10^{-5}, 10^{-4}, 10^{-3}, 0.01, 0.1, 1.0\}$.
- We apply every optimization algorithm to the Ridge Regression problem generated by each λ .

2. The Goals:

By establishing these different regularization classes, we can evaluate:

- **Regularization Impact:** How varying regularization conditions affects the convergence speed and performance of the *same* algorithm.
- **Inter-Algorithm Comparison:** How *different* algorithms perform against each other when subjected to the *same* regularization conditions.

Testing approach

3. Early Stopping Mechanism:

- At each iteration, we evaluate if the algorithm has successfully reached the minimum of the loss function:

$$\epsilon = f(w_T) - f(w^*) \leq 10^{-6}$$

- If convergence is achieved, execution terminates immediately, returning the elapsed time and total iterations.
- Otherwise, the algorithm continues until it hits a predefined `max_iters` limit.

4. Exact Analytic Solution (Benchmark):

- We have to compute the exact optimal solution for every Ridge Regression subproblem to use it as a benchmark.
- We use the closed-form solution by `np.linalg.solve(A, b)`.

λ	10^{-5}	10^{-4}	10^{-3}	0.01	0.1	1.0
$f(w^*)$	0.04265771	0.04276069	0.04369175	0.04977553	0.07101217	0.16355421

Gradient Descent - Naive Learning Rate

Gradient descent by definition updates the weight vector w step-by-step by moving in the opposite direction of the gradient, scaled by the learning rate (step size) γ :

$$w^{(t+1)} = w^{(t)} - \gamma \nabla f(w^{(t)})$$

For this first run, we are using a **naive, very small learning rate** $\gamma = 0.0001$.

Choosing a very small step size is a conservative and highly stable approach: it guarantees that our algorithm will not overshoot the minimum or diverge.

But this safety comes with a heavy trade-off in efficiency: taking such microscopic steps means the algorithm might require a massive number of iterations to finally reach the optimal solution.

Gradient Descent - Naive Learning Rate

Final results per λ using a naive learning rate ($\gamma = 0.0001$):

λ	Loss	γ	Iters	Time[s]
10^{-5}	0.05510	0.0001	50000	9.55
10^{-4}	0.05512	0.0001	50000	10.53
10^{-3}	0.05528	0.0001	50000	9.70
0.01	0.05689	0.0001	50000	10.24
0.1	0.07171	0.0001	50000	11.27
1.0	0.16356	0.0001	31209	6.40

Observations:

- For low regularizations (from $\lambda = 10^{-5}$ to 0.1), the algorithm hits the `max_iters` limit of 50.000 without reaching the tolerance of the early stop.
- Only for the strongest regularization ($\lambda = 1.0$), the algorithm is pushed fast enough to converge in exactly 31.209 iterations.

Gradient Descent - Assuming Bounded Gradient

Assume the distance between initial solution and optimal solution is bounded by R and this bounded region contains all iterates.

Since we initialize our initial solution vector to **Zero** (`np.zeros`), the distance coincides exactly with the norm of the solution:

$$\|w_0 - w^*\| \leq R$$

$$\|0 - w^*\| = \|w^*\| \leq R$$

The assumption on R means: *"The optimal solution and the optimization path to reach it lie within a sphere of radius R ."*

- Since we have **standardized the data** (X and y have zero mean and unit variance), the optimized weights w will naturally be small (typically between -1 and 1)
- We set $R = 5.0$. A generous safety margin (a "Bounded Region") that is virtually guaranteed to contain the true solution.

Gradient Descent - Assuming Bounded Gradient

We need to find a number B such that $\|\nabla f(w)\| \leq B$ for every w within the radius R . We rewrite the gradient formula in this way:

$$\nabla f(w) = \left(\frac{1}{n} X^T X + \lambda I \right) w - \frac{1}{n} X^T y$$

Using the triangle inequality ($\|a + (-b)\| \leq \|a\| + \|b\|$) and the definition of matrix norms:

$$\|\nabla f(w)\| \leq \underbrace{\left\| \frac{1}{n} X^T X + \lambda I \right\|}_{\text{This is smooth parameter } L} \cdot \underbrace{\|w\|}_{\leq R} + \underbrace{\left\| \frac{1}{n} X^T y \right\|}_{\text{Constant}}$$

So the formula for B is:

$$B = L \cdot R + \left\| \frac{1}{n} X^T y \right\|$$

Gradient Descent - Assuming Bounded Gradient

By making the assumption that our optimization path remains within a bounded region $\|w_0 - w^*\| \leq R$, and having calculated the maximum gradient bound B such that $\|\nabla f(w)\| \leq B$ for all w , we have effectively established that our objective function is **B -Lipschitz continuous** within this domain.

For Lipschitz continuous functions, optimization theory provides a specific, optimal step size: we can now set γ using the formula:

$$\gamma = \frac{R}{B\sqrt{T}}$$

where $T = 10.000$ iterations.

Gradient Descent - Assuming Bounded Gradient

Final results per λ using the bounded gradient step size (γ):

λ	Loss	γ	Iters	Time[s]
10^{-5}	0.04966	0.00214	10000	1.86
10^{-4}	0.04969	0.00214	10000	2.04
10^{-3}	0.04997	0.00214	10000	2.88
0.01	0.05263	0.00214	10000	2.15
0.1	0.07103	0.00210	10000	2.02
1.0	0.16356	0.00176	1768	0.32

Observations:

- **Larger Steps:** The dynamic $\gamma \approx 0.002$ is safely 20x larger than the naive guess.
- **Massive Speedup:** For $\lambda = 1.0$, it converges in only 1,768 iterations (vs. 31,209).
- **Higher Accuracy:** It reaches a strictly lower loss in just 10,000 steps compared to the 50,000 steps of the naive approach.

Gradient Descent - Smoothness

Previously we have seen that function is **L -Smooth**.

- **Value of L :** It is the largest eigenvalue of the Hessian matrix (the maximum curvature).

$$L = \lambda_{\max} \left(\frac{1}{n} X^T X \right) + \lambda$$

or equivalently

$$L = \left\| \frac{1}{n} X^T X \right\| + \lambda$$

- **Consequence:** This value is the **speed limit** for Gradient Descent. If we use a learning rate $\gamma > \frac{2}{L}$, the algorithm will diverge. The optimal safe step size is $1/L$.

Gradient Descent - Smoothness


Final results per λ using the optimal step size ($\gamma = 1/L$):

λ	Loss	$\gamma = 1/L$	Iters	Time[s]
10^{-5}	0.04266	0.23247	8353	1.56
10^{-4}	0.04276	0.23247	7767	1.74
10^{-3}	0.04369	0.23242	4438	1.76
0.01	0.04978	0.23193	765	0.22
0.1	0.07101	0.22719	142	0.07
1.0	0.16355	0.18862	15	0.01

Observations:

- **Optimal Steps:** The safe learning rate ($\gamma \approx 0.23$) is $\sim 2,300\times$ larger than our naive guess.
- **Lightning Fast:** For $\lambda = 1.0$, it perfectly converges in just 15 iterations (down from 31.209).
- **Perfect convergence:** For all case the algorithm converge to optimal solution.

Convergence Analysis: Comparing Gradient Descent implementations



Nesterov's Accelerated Gradient Descent

Now we implement the **Accelerated Gradient Descent** algorithm. This method improves standard gradient descent by computing the next iteration point as a dynamically weighted average between a conservative "smooth" step and a more "aggressive" step.

Choosing arbitrary initial points $\mathbf{z}_0 = \mathbf{u}_0 = \mathbf{w}_0$, for $t \geq 0$ we set:

$$\begin{aligned}\mathbf{u}_{t+1} &:= \mathbf{w}_t - \frac{1}{L} \nabla f(\mathbf{w}_t) \\ \mathbf{z}_{t+1} &:= \mathbf{z}_t - \frac{t+1}{2L} \nabla f(\mathbf{w}_t) \\ \mathbf{w}_{t+1} &:= \frac{t+1}{t+3} \mathbf{u}_{t+1} + \frac{2}{t+3} \mathbf{z}_{t+1}\end{aligned}$$

- $\mathbf{w}$ represents the main weight vector.
- $\mathbf{u}$ represents the prudent/smooth update.
- $\mathbf{z}$ represents the aggressive update.

Nesterov's Accelerated Gradient Descent

Final results per λ using Accelerated Gradient Descent:

λ	Loss	Iters	Time[s]
10^{-5}	0.04266	203	0.10
10^{-4}	0.04276	202	0.12
10^{-3}	0.04369	152	0.07
0.01	0.04978	86	0.05
0.1	0.07101	23	0.01
1.0	0.16355	12	0.005

Observations:

- **Massive Speedup:** For ill-conditioned problem ($\lambda = 10^{-5}$), iterations go down from 8,353 to just 203 ($\sim 40\times$ improvement).
- **Faster Convergence Rate:** Validates the theoretical improvement from $O(1/\epsilon)$ to $O(1/\sqrt{\epsilon})$.
- **Diminishing Returns:** For $\lambda = 1.0$ (a steep, well-conditioned bowl), the advantage is marginal (12 vs. 15 iterations).

Stochastic Gradient Descent with Mini-batches

In standard Gradient Descent, computing the exact gradient requires evaluating the entire dataset (all n samples) at every single step. While highly accurate, this becomes computationally prohibitive and slow for massive datasets. **Mini-batch Stochastic Gradient Descent (SGD)** addresses this bottleneck by approximating the true gradient using a much smaller, randomly selected subset of the data of size m (where $m < n$).

Starting from an initial point $\mathbf{x}_0 \in \mathbb{R}^d$, at each iteration t we choose a random subset of indices $\{h_1, \dots, h_m\} \subseteq \{1, \dots, N\}$ and update the weights using a learning rate $\gamma_t \geq 0$:

$$\mathbf{w}_{t+1} := \mathbf{w}_t - \gamma_t \frac{1}{m} \sum_{j=1}^m \nabla f_{h_j}(\mathbf{w}_t)$$

Stochastic Gradient Descent with Mini-batches

To realize the random subset selection, the algorithm operates in "epochs". At the start of an epoch, the entire dataset is randomly shuffled (`np.random.permutation`). Within an epoch, a loop then sequentially iterates over this shuffled data in chunks of size `batch_size` (our m).

The `minibatch_gradient` function then calculates the scaled gradient on this small chunk, and the main loop updates the weights.

Just as we did in our initial experiments with standard Gradient Descent, we will establish a baseline by using a simple, "naive" constant learning rate of $\gamma = 10^{-4}$, with a `batch_size = 1000`.

Stochastic Gradient Descent with Mini-batches

Final results per λ using Stochastic Gradient Descent with a naive learning rate ($\gamma = 0.0001$):

λ	Loss	γ	Iters	Time[s]
10^{-5}	0.05510	0.0001	50000	12.79
10^{-4}	0.05512	0.0001	50000	13.10
10^{-3}	0.05528	0.0001	50000	13.14
0.01	0.05689	0.0001	50000	12.84
0.1	0.07171	0.0001	50000	13.54
1.0	0.16356	0.0001	31254	7.35

Observations:

- Results are almost identical to our initial Naive Full-Batch Gradient Descent. The algorithm still hits the hard limit of 50,000 iterations for small λ values, getting stuck at a suboptimal loss of roughly 0.05510 instead of reaching the true minimum (≈ 0.04266).

Mini batches SGD - Decaying Learning Rate

Now try doing a mini-batch SGD with a better learning rate using the fact that the objective function is μ -strongly convex. From **Tame Strong Convexity**, the learning rate to use is:

$$\gamma_t = \frac{2}{\mu(t+1)}$$

and it decays over each iteration.

However, applying this formula naively when our regularization parameter λ is small leads to immediate divergence.

- If μ is very small, the initial step size γ_t becomes very large.
- This massively exceeds our theoretical L -smoothness "speed limit" ($2/L$), causing the algorithm to diverge.

Mini batches SGD - Decaying Learning Rate

Looking closely at our calculated strong convexity parameters (μ), we can demonstrate why applying the classic theoretical decay rate $\gamma_t = \frac{2}{\mu(t+1)}$ causes divergence.

For our smallest ridge penalization ($\lambda = 10^{-5}$), the minimum curvature of our loss function is extremely flat, resulting in $\mu \approx 0.00130$. If we plug this value into the pure **Tame Strong Convexity** formula to calculate the learning rate for our very first iteration ($t = 0$), we get:

$$\gamma_0 = \frac{2}{0.00130} \approx 1538$$

This initial step size is catastrophically large. As we established earlier, the absolute maximum safe step size before an algorithm diverges is bounded by the smoothness parameter $(2/L) \approx 0.46$. An initial step of 1538 massively violates this theoretical "speed limit".

This same issue occurs across all other small values of the ridge regularization parameter λ .

Mini batches SGD - Decaying Learning Rate

To safely apply the optimal decay rate by the **Tame Strong Convexity**, we propose an adjusted learning rate:

$$\gamma_t = \frac{2}{L + \mu(t + 1)}$$

In this way we get:

- **Initial Stability ($t = 0$):** At the very first iteration, the learning rate evaluates to exactly $\frac{2}{L+\mu} < \frac{2}{L}$. This respects the theoretical optimal speed limit for smooth strongly convex functions, acting as a safety cap that prevents the algorithm from diverging, regardless of how small μ is.
- **Asymptotic Convergence ($t \rightarrow \infty$):** As the iterations progress, the L term becomes dominated by μt . The learning rate transitions to approximate $\frac{2}{\mu t}$, strictly adhering to the decay rate required to suppress the stochastic noise of the mini-batches and guarantee convergence to the global minimum.

Mini batches SGD - Decaying Learning Rate

Final results per λ using the adjusted decaying learning rate (γ_t):

λ	Loss	Iters	Time[s]
10^{-5}	0.04266	14199	3.58
10^{-4}	0.04276	15115	3.94
10^{-3}	0.04369	8104	2.23
0.01	0.04978	1730	0.71
0.1	0.07101	404	0.13
1.0	0.16355	101	0.04

Observations:

- **Guaranteed Convergence:** Safely bounded initial steps eliminate chaotic bouncing, achieving perfect convergence to the exact global minimum for every λ .
- **Noise Reduction:** To settle exactly at the minimum, the decaying learning rate must become extremely small, forcing thousands of tiny, cautious steps to average out mini-batch noise.

Newton's Method

Until now, we have relied on first-order optimization algorithms (like Gradient Descent and its variants), which only use the gradient (the slope) to determine the direction of our next step. **Newton's Method** introduces a massive upgrade by incorporating second-order derivative information—specifically, the curvature of the loss landscape represented by the Hessian matrix $\nabla^2 f(\mathbf{x})$.

By using both the slope and the curvature, the algorithm doesn't just step blindly down the hill; it models the exact shape of the convex bowl and points directly toward the bottom.

Mathematically, for minimizing a convex function $f : \mathbb{R}^d \rightarrow \mathbb{R}$ in the d -dimensional case, the update rule is defined as:

$$\mathbf{w}_{t+1} := \mathbf{w}_t - \nabla^2 f(\mathbf{w}_t)^{-1} \nabla f(\mathbf{w}_t)$$

Newton's Method

Final results per λ using Newton's Method:

λ	Loss	Iters	Time[s]
10^{-5}	0.04266	1	0.006
10^{-4}	0.04276	1	0.004
10^{-3}	0.04369	1	0.003
0.01	0.04978	1	0.003
0.1	0.07101	1	0.003
1.0	0.16355	1	0.002

Observations:

- **Single update iteration:** Newton's Method achieves exact numerical convergence (to our 10^{-6} tolerance) in a single update step.

Newton's Method

This result is a direct consequence of the underlying mathematical nature of **Ridge Regression**. The objective function is a quadratic and strictly convex function with respect to the weight vector w .

$$f(w) = \frac{1}{2n} \|Xw - y\|^2 + \frac{\lambda}{2} \|w\|^2$$

The One-Step Convergence Lemma:

On nondegenerate quadratic functions, with any arbitrary starting point $w_0 \in \mathbb{R}^d$, Newton's method yields the exact global minimum w^ in exactly one step ($w_1 = w^*$).*

Because the Ridge Regression objective function satisfies the condition of a nondegenerate quadratic function (a matrix with strictly positive eigenvalues has a non-zero determinant and is therefore invertible), this experimental behavior aligns exactly with optimization theory.

Adagrad

Adagrad modifies standard Stochastic Gradient Descent by introducing a dynamic, per-coordinate learning rate. At each iteration t , we first pick a stochastic gradient $\mathbf{g}_t$.

1. Gradient Accumulation

We accumulate the history of squared gradients for each individual coordinate i :

$$[G_t]_i := \sum_{s=0}^t ([\mathbf{g}_s]_i)^2 \quad \forall i$$

(This step tracks how strong the gradients have been in the past for the i -th coordinate).

2. Parameter Update

We then update each coordinate i of our parameter vector $\mathbf{w}$:

$$[\mathbf{w}_{t+1}]_i := [\mathbf{w}_t]_i - \frac{\gamma}{\sqrt{[G_t]_i}} [\mathbf{g}_t]_i \quad \forall i$$

The Adaptive Learning Rate:

The crucial part of this formula is the term $\frac{\gamma}{\sqrt{[G_t]_i}}$, which acts as our new coordinate-specific learning rate.

- By dividing the base learning rate γ by the square root of the accumulated gradients, the algorithm automatically **reduces the step size** for coordinates that have historically had large gradients.
- Instead, if a coordinate has had consistently small gradients, its learning rate remains relatively **higher**.

This mechanism speeds up convergence and reduces oscillations along directions where gradients are already large.

Adagrad

Final results per λ using the Adagrad algorithm with $\gamma = 0.1$:

λ	Loss	Iters	Time[s]
10^{-5}	0.04266	42511	11.99
10^{-4}	0.04276	37110	9.61
10^{-3}	0.04369	19376	4.73
0.01	0.04978	3477	0.78
0.1	0.07101	9059	1.98
1.0	0.16356	8292	1.92

Observations:

- **Guaranteed Convergence:** Adagrad successfully reached the exact global minimum for all λ values.
- **High Curvature Penalty:** Strong regularization (e.g., $\lambda = 1.0$) creates a steeper convex bowl. This produces massive gradients in the very first iterations, which permanently push up the denominator G_t . The algorithm push down γ too early, requiring thousands of tiny steps to cross the final tolerance threshold.

Summary of Results for $\lambda = 10^{-5}$

Performance comparison on the ill-conditioned problem ($\lambda = 10^{-5}$):

Optimization Method	Loss	Iterations	Time[s]
Naive Gradient Descent	0.05510	50000*	~10.50
Bounded Gradient Descent	0.04966	10000*	1.86
Smooth Gradient Descent	0.04266	8353	1.56
Accelerated Gradient Descent	0.04266	203	0.10
Naive SGD	0.05510	50000*	12.79
Decaying Learning Rate SGD	0.04266	14199	3.58
Newton's Method	0.04266	1	0.006
Adagrad	0.04266	42511	11.99

**Hit maximum iterations limit without converging*

Summary of Results for $\lambda = 1.0$

Performance comparison on the strongly regularized problem ($\lambda = 1.0$):

Optimization Method	Loss	Iterations	Time[s]
Naive Gradient Descent	0.16356	31209	6.39
Bounded Gradient Descent	0.16356	1768	0.32
Smooth Gradient Descent	0.16355	15	0.01
Accelerated Gradient Descent	0.16355	12	0.005
Naive SGD	0.16356	31254	7.35
Decaying Learning Rate SGD	0.16355	101	0.04
Newton's Method	0.16355	1	0.002
Adagrad	0.16356	8292	1.92

Conclusions

- **Impact of λ on Convergence:**

Increasing λ helps algorithms to converge, making the optimization's problem better conditioned. A higher λ drastically reduces iterations and execution time.

- **Algorithm Performance Comparison:**

- **Newton's Method:** Converges in exactly 1 step across all λ by leveraging second-order curvature.
- **Accelerated GD:** Best first-order method, can effortlessly navigate flat valleys.
- **Standard GD ($1/L$):** Highly efficient for well-conditioned (high- λ) problems, but struggles heavily in flat scenarios.
- **Mini-batch SGD:** Safely converges, but stochastic noise makes strict high-precision possible only when learning rate becomes very small.
- **Adagrad:** Auto-tunes step sizes without theoretical parameters, but the monotonically shrinking learning rate causes slow convergence.